

## Image Convolution on FPGAs:

### the Implementation of a Multi-FPGA FIFO Structure

Arrigo Benedetti<sup>\*†</sup>, Andrea Prati<sup>\*</sup>, Nello Scarabottolo<sup>\*</sup>

- <sup>\*</sup> Dept. of Engineering Sciences, Università di Modena, via Campi, 213, I41100 Modena, Italy  
phone +39 59 376739, fax +39 59 376799, e-mail {benedett, prati, scarabot}@dsi.unimo.it  
<sup>†</sup> presently at Computational Vision Laboratory of California Institute of Technology, Pasadena

#### Abstract

In this paper, we present an implementation of a real-time convolver, based on Field Programmable Gate Arrays (FPGA's) to perform the convolution operations. Main characteristics of the proposed approach are the usage of external memory to implement a FIFO buffer where incoming pixels are stored and the partitioning of the convolution matrix among several FPGA's, in order to allow data-parallel computation and to increase the size of the convolution kernel.

#### 1. Introduction

The work presented in this work is one of the result of a research activity aimed at implementing dedicated architectures for image processing. In particular, the research activity focuses on reconfigurable devices like Field Programmable Gate Arrays (FPGA's) [1] as the most promising ones in applications where processing speed – typical of dedicated solutions – has to be matched with low-cost, flexible systems capable of performing several different tasks.

Among the various algorithms frequently encountered in image processing, we put our attention to 2-D convolution, where a  $M \times N$  pixels input image has to be convoluted with a  $K \times R$  kernel to obtain an output image where each pixel depends on a  $K \times R$  window of neighboring pixels in the input image [2] [3], as shown in Fig. 1.

Even if the problem is already well known and solved both in software and in dedicated hardware, a satisfactory solution for FPGA-based real-time convolution is at our knowledge still lacking.

To reach such a real-time behavior, it is in fact necessary to process input image pixels – entering the convolution system in the usual raster scan order row-by-row – to ensure an equal rate of output pixels (to avoid unlimited buffering capabilities) and to minimize




Figure 1. Convolution of a raster-scan image  
( $K=R=3$ ,  $M=N=5$ )

processing latency.

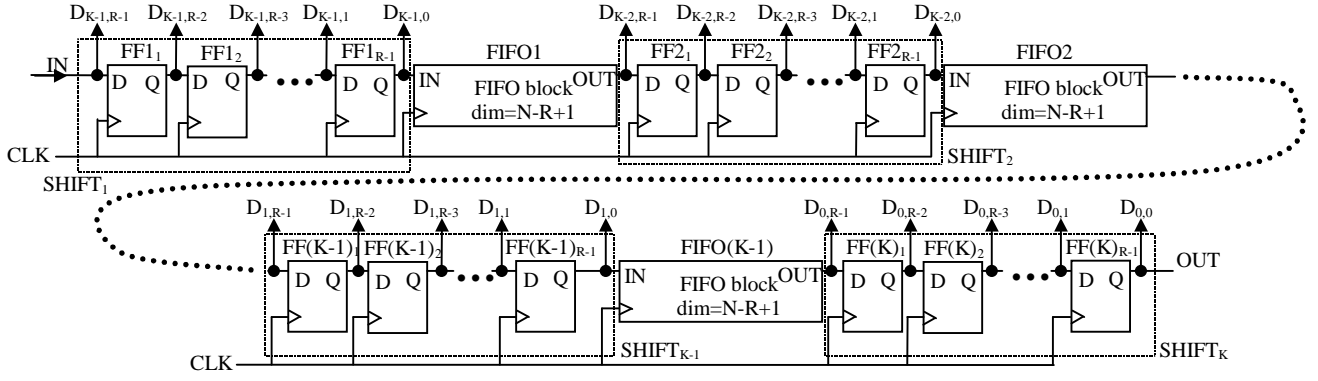
In this perspective, a software based solution, where the overall image is stored in RAM memory before starting FPGA computation, seems inadequate for the added latency and the necessity of a large number of memory accesses, possibly exceeding the real-time constraints.

A far better approach consists in the implementation of a pipeline processor, where pixels enter the pipeline and produce results as soon as the convolution window has been filled. Referring again to Fig. 1, this means that we need to store  $(K-1) \times N + R$  pixels and a processing latency – relative to the central pixel of the convolution window – given by

$$Lat = \left\lceil \frac{R}{2} \right\rceil + N \times \left\lceil \frac{K}{2} \right\rceil$$

is introduced.

To accomplish this, it is possible to adopt a set of shift registers [4], implementing a FIFO memory able to store the  $(K-1) \times N + R$  pixels above mentioned. For a  $n$  bit per pixel image (i.e., a  $2^n$  colors/gray-levels image) this leads itself to  $n$  shift registers as shown in Fig. 2, where single flip-flops store pixels belonging to the convolution window (i.e., pixels required to perform the current convolution product) and FIFO blocks store pixels already grabbed and necessary for subsequent convolution



**Figure 2. Memory structure allowing real-time convolution**

products.

This approach – quite obvious in dedicated devices – is not well suited to an FPGA implementation. In fact, commercially available FPGA devices contain a limited number of memory elements, since they are constituted by few hundreds of Configurable Logic Blocks (CLB's), each provided with one or two flip-flops and a configurable logic network. Implementing a memory structure on FPGA's means using a lot of devices, thus wasting their (re)configuration properties (e.g., with rather large FPGA devices like the Xilinx XC4010E, a 3×3 convolution matrix applied to a 512×512 image using 8 bits per pixels needs  $(512 \times 2 + 3) \times 8 = 1027 \times 8 = 8216$  flip-flops, requiring 11 FPGAs).

In fact, previous works (as [5] and [6]) have presented methods for synthesizing FIFO memories using the flip-flops present inside the CLB's of the Xilinx XC4000 series devices: unfortunately, the number of available CLB's permits the implementation of FIFO memories far too small for our needs. For this reason, the solution we propose implements the shift registers of Fig. 2 using also external RAM memory interfaced with the convolution FPGA's. In particular, flip-flops storing pixel values belonging to the convolution window are stored into internal CLB's, while FIFO blocks storing pixels necessary for subsequent convolution operations are stored into the external RAM memory.


The implementation of such a FIFO and its possibilities in terms of speed and kernel size are the main topics of this paper.

In Sec. 2, we discuss the requirements of the FIFO memory and the solution that we have implemented for interfacing it with a single FPGA.

In Sec. 3, we extend FIFO capabilities to allow parallel processing by several FPGAs, thus relaxing constraints on the convolution matrix dimension.

In Sec. 4, an experimental prototype written in VHDL language and implemented in a G800 GigaOps Spectrum board containing 8 Xilinx XC4010E devices is discussed.

Finally, possible developments are outlined in Sec. 5.



**Figure 3. Subsequent steps in the convolution of a raster-scan image ( $K=R=3$ ,  $M=N=5$ )**

## 2. The FIFO memory interfaced with a single FPGA

To better understand the requirements of the external FIFO memory, it is worth to refer to Fig. 3.

In order to minimize the number of external memory accesses necessary to obtain the required FIFO behavior, we propose an approach based on the following assumptions:

- each new pixel entering the convolution system ( $D_{2,3}$  in Fig. 3-b) is routed both to the FPGA convolver and to the external FIFO memory, where it is stored in a single write cycle;
- the  $K-1$  pixels belonging to the convolution “column” of the new pixel ( $D_{0,3}$  and  $D_{1,3}$  in Figure 3-b) must be read in  $K-1$  read cycles in order to be routed to the FPGA convolver;
- the remaining  $(R-1) \times K$  pixels ( $D_{0,1}$ ,  $D_{0,2}$ ,  $D_{1,1}$ ,  $D_{1,2}$ ,  $D_{2,1}$  and  $D_{2,2}$  in Figure 3-b) necessary to perform the convolution are stored inside the FPGA.

Summarizing, for each new pixel coming from the video signal source via the frame grabber, 1 memory write and  $K-1$  memory reads must be performed to correctly feed the FPGA convolver.




Figure 4. Behavior of the external FIFO ( $N=M=5$ ,  $K=R=3$ )

This has been accomplished in the FPGA system by implementing the FIFO memory in a circular buffer stored in RAM, based on the following elements:

1. The input/output clock signal, here called *video clock*, which synchronizes frame grabber and video encoder operations. During each cycle of clock, a new pixel enters the convolver and a new complete convolution takes place.
2. A *convolution clock*, whose frequency  $f_{conv}$  is  $K$  times higher than the frequency  $f_{play}$  of the play clock. During each cycle of the convolution clock, a memory access (either read or write) to the external FIFO is performed.
3. A write pointer  $counter_{wr}$ , stored inside the FPGA convolver is used to address the first free location inside the external FIFO, where the incoming pixel has to be stored.
4.  $K-1$  read pointers  $counter_{rd(i)}$ , also implemented inside the FPGA are used to address the locations of the pixels to be read from the external FIFO to complete the computation of the convolution window.

Using the above elements, the external FIFO, requiring a memory able of accommodating at least  $(K-2) \times N + (N-R+1) + 1 = (K-1) \times N - R + 2$  pixels, is obtained

as summarized in Fig. 4, where a  $3 \times 3$  convolution window and a  $5 \times 5$  input image are assumed.

Before starting to fill the FIFO,  $R-1$  pixels have to be shifted into the  $SHIFT_1$  register (Fig. 4.(a)). The next  $N-R+1$  pixels are stored in the FIFO denoted FIFO1 in Fig. 2, filling it until  $counter_{wr}$  reaches location 0002 (Fig. 4.(b)). At this time  $counter_{rd(1)}$  is enabled and starts to address data from FIFO1 and to shift them into  $SHIFT_2$ . After  $N$  additional pixels (i.e., a full image row has been read)  $counter_{rd(0)}$  is enabled in turn and begins to store data in  $SHIFT_3$ . The situation is shown in Fig. 4.(c). After  $N-R+1$  additional pixels, the first convolution window is available and the FPGA starts the computation using  $counter_{rd(0)}$  and  $counter_{rd(1)}$  to fetch  $D_{02}$  and  $D_{12}$  from the FIFO (Fig. 4.(d)).

It should be noticed that this approach poses the following limits on the dimension of the convolution matrix:

- 1) a first limit derives from the maximum frequency allowed for the convolution clock, used to access the external RAM memory to fetch the next "column" of pixels. This limit affects then the maximum number of rows (thus the value of  $K$ ) in the convolution matrix: referring to the 8.33 MHz of the video clock (adopted by the european PAL video standard) and

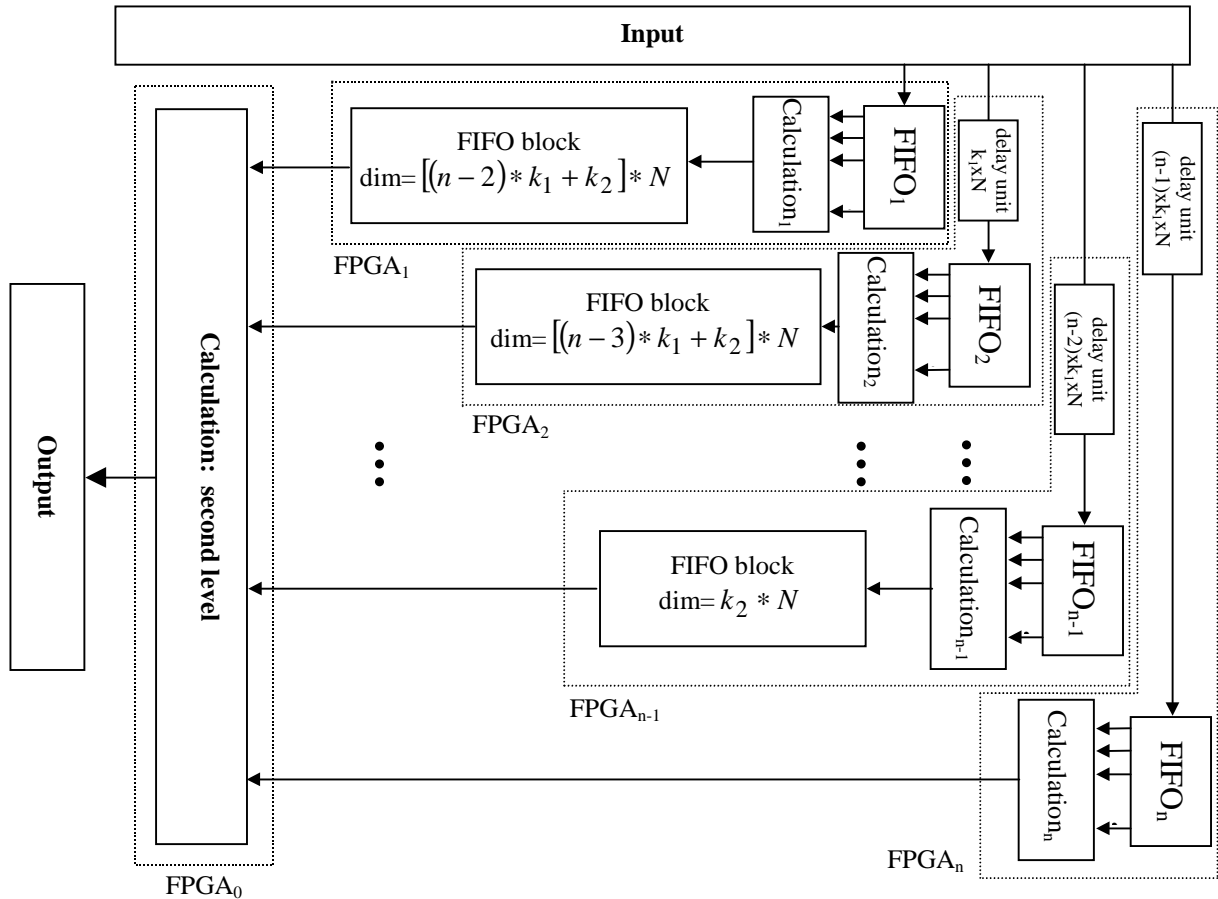


Figure 5. Multi-FPGA architectural schema

using 20 nsec access time RAM (i.e., 50 MHz maximum  $f_{conv}$ ), this results in a convolution matrix of at most 6 rows;

- 2) a second limit derives from the capacity of the FPGA convolver, that, even in case of particularly favorable convolution matrices (containing only 0s, 1s and power-of-2 coefficients, thus requiring only logical operations and shifts - e.g. Sobel, Prewitt and Laplace matrices [2] - which can be implemented using few CLBs), has to store inside the internal flip-flops the  $K \times R$  pixels to be handled. This limit affects then the overall dimension of the convolution matrix.

It is easier shown that the most critical parameter in terms of capability of handling large convolution kernels is the number  $K$  of rows; moreover, it can be noticed that, for a given "area" of the convolution matrix, it is more convenient to minimize  $K$  and maximize  $R$  in order to relax clock constraints.

On the basis of the above considerations, we have devised a multi-FPGA structure that is presented in the next section.

### 3. Multi-FPGA solution

When several FPGA devices are available, the most convenient way of partitioning computation among them is represented in Fig. 5:  $n$  FPGAs are used to process in parallel  $n$  different portions of the convolution window, and an additional FPGA (the 0-th one in Figure 5) merges intermediate results performing (usually) simple sum operations. This would mean to subdivide the  $K \times R$  convolution matrix into  $n$  sub-matrices of size  $k \times r$ , where the following relations hold:

$$k = \frac{K}{p}, \quad r = \frac{R}{q}, \quad p \times q = n. \quad (1)$$

A great deal of work exist about partitioning FPGA designs (e.g., [7], [8] and [9]), however, given the considerations expressed in previous section, the best choice for the convolution matrix subdivision in our case is given by

$$p = n, \quad q = 1,$$

which means a set of horizontal "stripes".

It must be noticed however that the second level FIFO blocks shown in Fig. 5 have to be included to ensure that the results of each single convolution stripe reach FPGA<sub>0</sub> aligned in time. These FIFO's introduce the need of two additional accesses to external RAM: one to write the input and one to read the output. In other words, all the FPGA's but the last one must use two cycles of the convolution clock to access the second level FIFO's, which implies two less accesses to their first-level FIFO, which contains the pixel values to be multiplied.

Since the number of allowed accesses to first-level FIFO's is the upper limit to the number of rows in each stripe, and since the last FPGA does not need a second level FIFO, the best subdivision of the convolution matrix

can be obtained as follows.

Referring to Fig. 5, we see that the original  $K$  rows must be partitioned into  $p$  stripes of  $R$  columns: one stripe of  $k_2$  rows and  $p-1$  stripes of  $k_1$  rows. For previous considerations,  $k_2 = k_1 + 2$ . Obviously, to correctly partition the original  $K$  rows,  $k_1 \times (p-1) + k_2 = K$  must hold. From this equation we can show that

$$k_1 = \frac{K-2}{n}, \quad k_2 = \frac{K-2}{n} + 2. \quad (2)$$

Delay units in Fig. 5 allow each FPGA to obtain its own stripe of data.

Besides the advantages already discussed, this approach allows to optimally map the convolution matrix into  $n$  FPGAs provided that:

$$(K-2) \bmod n = 0$$

without any other constraint upon  $R$ .

To estimate the effectiveness of such approach, let us revise the practical example made at the end of previous section. For a 20 nsec access time RAM to be used in a system working with the european PAL video standard, we have:

- $f_{play} = 8.33$  MHz
- maximum  $f_{conv} = 50$  MHz
- maximum number  $k$  of accesses to external RAM in a video clock cycle:  $\left\lfloor \frac{50}{8.33} \right\rfloor = 6$
- maximum  $K$  dimension of the convolution window:  $(n-1) \times (6-2) + 6 = 4n + 2$

For a system with  $n=8$  FPGAs, using only square kernels, this means a  $34 \times 34$  upper limit to the kernel dimension.

If we remove the choice of square matrices, it is possible to further increase  $R$  up to the space limit given by the number of available CLB's.

### 4. Prototypal implementation

In this section, we present a VHDL [10] implementation of the multi-FPGA structure, discussed in previous section, on a multi-FPGA board designed for rapid prototyping [11].

To this purpose, we first describe the main characteristics of the prototyping board, with particular reference to the ones limiting the degrees of freedom in mapping computation and data paths onto the available FPGA's. Then, we show our solution and we evaluate the results.

#### 4.1 Prototypal board

The prototyping board we used is the GigaOps G800 Spectrum board [12], schematically shown in Fig. 6.

In Fig. 6 it is possible to notice the main blocks of this system.

- The actual computation is performed by pairs of Xilinx XC4010E FPGA's, connected in modules




Figure 6. Block diagram of the GigaOps G800 board

called *XMOD*'s: in Fig. 6, four modules (*MOD0* thru *MOD3*) are shown. These FPGAs are named *YPGA* and *XPGA* (from the name of the bus they are connected with). Both these FPGAs have two memory ports: one connected only to a 2 MBytes DRAM and one connected both to a 2 MBytes DRAM and to a 128 KBytes SRAM device. *XPGA* and *YPGA* communicate through a bus switch on the first memory port. This switch works on two virtual busses: a 16-bit data bus and a 10-bit address bus. It is important to stress that only *YPGA*'s are connected to *YBUS*, i.e. to the input/output data bus.

- A module called *SCVIDMOD* (S-VIDEO, COMPOSITE, VIDEO MODULE), that decodes/encodes video signals (PAL or NTSC). This module interfaces to *YBUS* for data input and output.
- An input FPGA (here called *VLPGA*) connected to the VESA local bus of the PC hosting the board. The *VLPGA* is interfaced with the *HBUS* and the *YBUS*. It contains all the registers needed for correct board operation (e.g., the *CLKMODE* register, that sets frequencies of the clocks distributed on the board).
- An output FPGA (here called *VMC*) connected to *SCVIDMOD*. This is an additional FPGA, directly interfaced with the video output and the *XBUS*.
- Three main busses that allow connections among the various blocks of the board. These busses are:
  - *YBUS*, a 32-bit I/O bus connected with *VLPGA*, *VMC* and the *YPGA*'s of the *XMOD*'s.
  - *HBUS*, a 16-bit bus used to configure and to load the FPGA's.
  - *XBUS*, a 64-bit bus normally used as four 16-bit data busses. Each of these busses is connected only to the *XPGAs* of the and to *VMC*.

The main data path of our application is the following: pixels generated by the video decoder are passed through the *YBUS* both to *VLPGA* and to the *YPGA*'s of the *XMOD*'s. These modules process data and pass the results either to the *VMC FPGA* through *YBUS* or to the *XPGA*'s through the bus switches. In the latter case, the *XPGA*'s can perform a further computation or simply pass the results to the *VMC FPGA* through the 64-bit *XBUS*. In

both cases, the *VMC FPGA* outputs the results of its processing upon data coming from the *XBUS* or *YBUS*.

## 4.2 The convolver prototype

A first characteristic that limits the degrees of freedom of this system is the kind of access to input/output data. In fact, *XPGA*'s are not connected to the *YBUS*, therefore they receive data to be processed only through the *YPGA*'s. This affects the effectiveness of the implementation: in particular, it makes more difficult to handle a multi-FPGA architecture that uses also *XPGAs* than an architecture using only *YPGAs*.

Another limit is constituted by the busses that interconnect *XPGA*'s with *YPGA*'s. These busses are also used to access to the 2 MBytes DRAM's. It is then extremely difficult to manage and to synchronize both connections with *XPGA* and with DRAM.

It is easy to note that the above restrictions make very hard and not effective to use both *XPGAs* and *YPGAs* to map computations.

On the contrary, *XPGAs* become useful to implement the second level FIFOs of Figure 5: this approach allows, in fact, to relax the constraints on the maximum number of accesses defined in the previous section. Thus, we can perform on each *YPGA*'s a convolution with a  $k \times r$  kernel.

Referring to the above considerations we decided to implement our prototype using only *YPGA*'s to perform convolution. In particular, since our own device has four *XMOD*'s, our multi-FPGA convolver uses the 4 *YPGAs* to map computation and the 4 *XPGA*'s to simply implement second level FIFO's.

Some additional considerations are needed to completely understand our prototype. Since every memory access requires first to present the memory address and to assert control signals (like memory read and memory write) then to deactivate these control signals before the next memory access, it is necessary to identify two timing events for starting and ending each memory access.

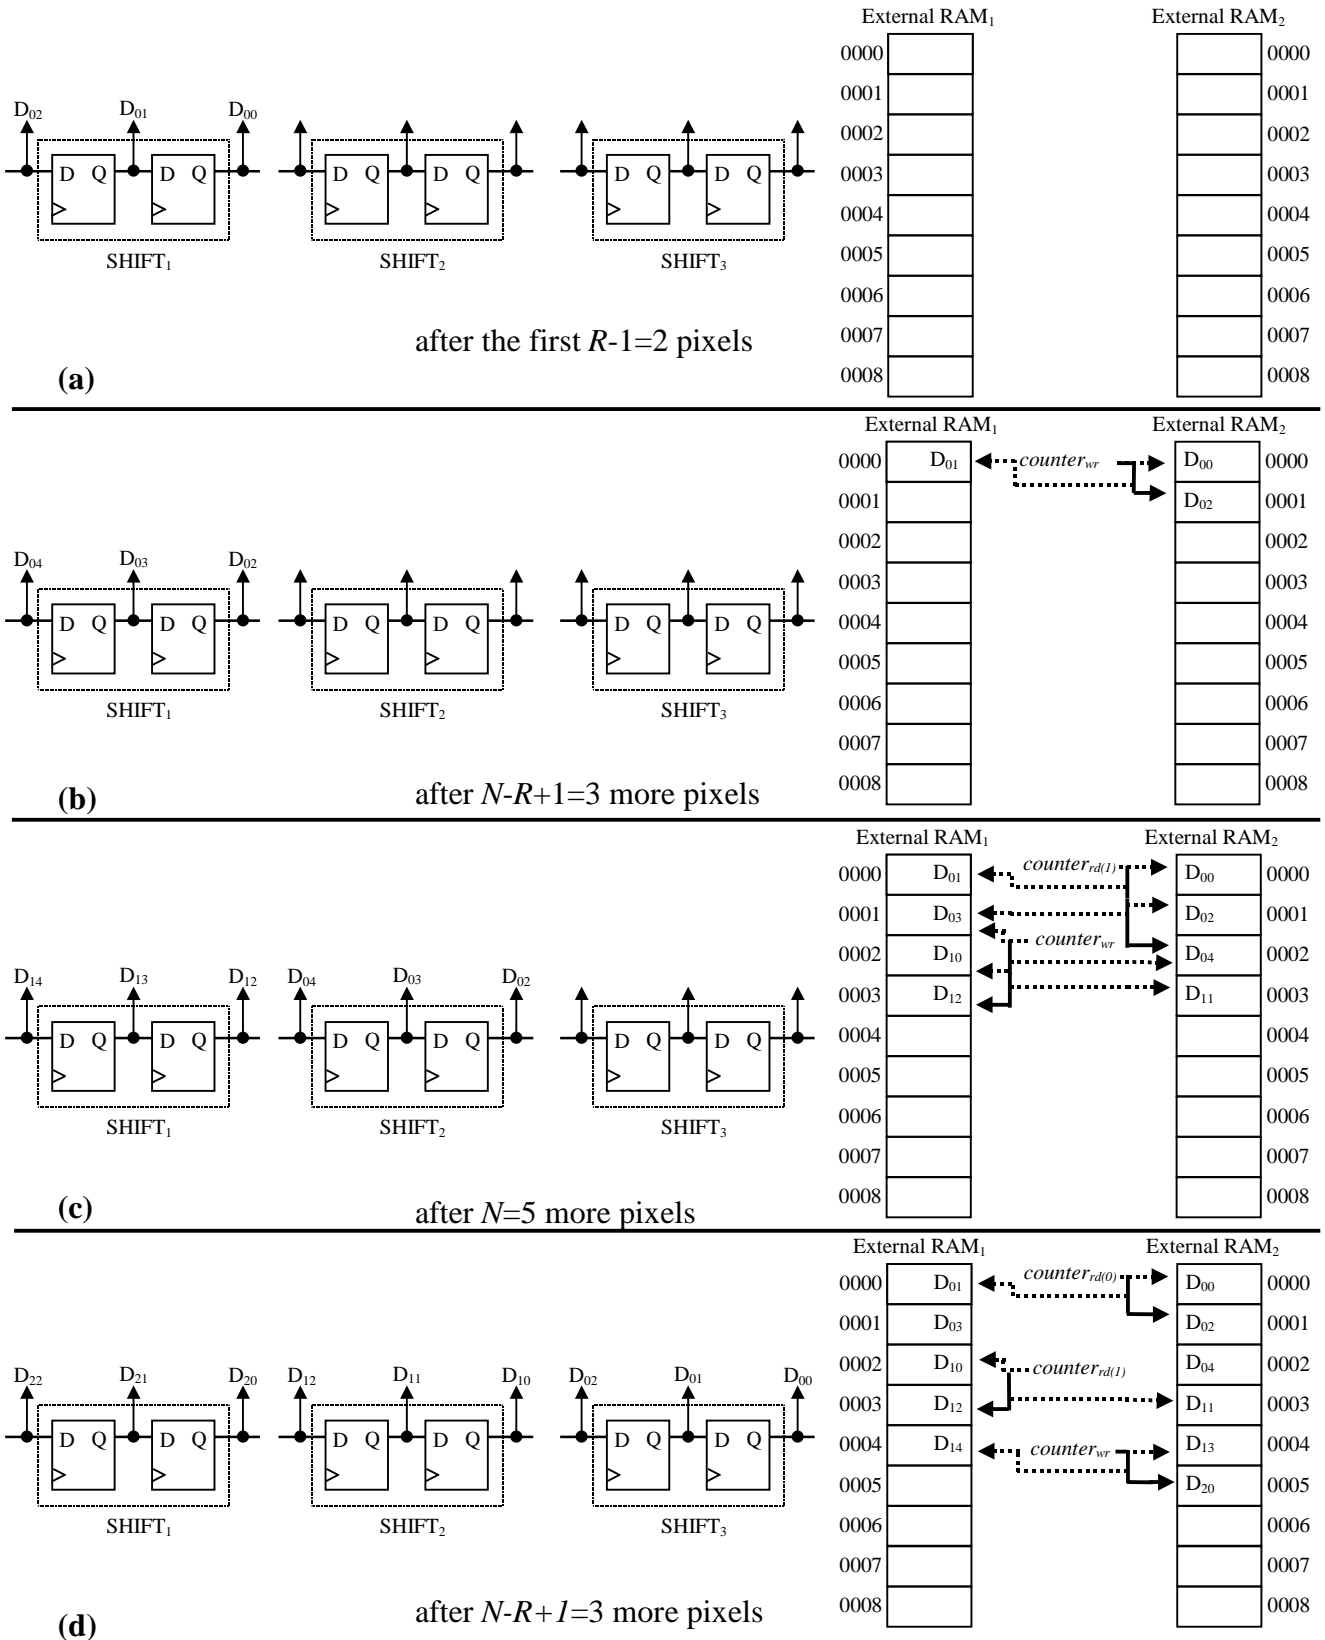


Figure 7. Behavior of the external FIFO implemented on dual-port memory ( $K=R=3$ ,  $N=M=5$ )

Due to board and memory characteristics, this poses an additional limitation: the maximum-frequency clock available (33.33 MHz: four times the PAL video standard  $f_{play}$ ) generates an event (either a rising edge or a falling edge) every 15 nsecs. This does not match the 30 nsec access time EDO DRAM access time (present on recent versions of the *G800* board), nor the 60 nsec access time DRAM (present on our prototype board), thus implying to reduce the performance by using only some edges of the clock waveform.

To avoid this problem, strictly related to the type of board and memory chips available, we exploited the second RAM port present on the *X210* modules, to perform two memory accesses at a time. Clearly, we must properly alternate memory writes and reads on the two ports. In Fig. 7 an example is reported.

With this approach, the FPGA's make two memory accesses per cycle of video clock to each of the two RAM banks available, allowing to reach an upper limit of a 16×16 convolution matrix with 30 nsec access time EDO DRAM and an upper limit of an 8×8 kernel with a 60 nsec access time DRAM. In fact, in this case, FPGA's can perform only one memory access per cycle of play clock to each of the two banks.

## 5. Conclusion and future work

In this paper, we have presented the architecture of a real-time convolver implemented by using FPGA's. In particular we have focused on the constraints that limit the maximum matrix dimension, which is directly related to the effectiveness of convolver. We have described a possible solution to relax these constraints, based on data parallelism.

In our prototype, we started testing the multi-FPGA architecture with several 3×3 convolution matrices, e.g., Sobel, Prewitt, Kirsch, Laplace and Cantoni operators [2] [3]. The limited dimension of these convolution matrices allowed us to perform convolution row-by-row on three different *YPGA*'s, without requiring first-level FIFO memories, thus avoiding the two external memory accesses discussed in the previous section.

Second level FIFO's and final calculations are performed by the *XPGA*'s and the VMC, respectively.

This choice led to meet the real-time constraints also in our *G800* prototypal board: in fact, even with a 60 nsec access time DRAM, the dual port used by the *XPGAs* allows to implement second-level FIFOs.

The next step will be the implementation of an 8×8 convolution matrix, spreading over 4 *YPGAs* the 8 convolution rows and using again the *XPGAs* for second level FIFO's. The fact that each *YPGA* has to process two convolution rows implies the implementation of first level FIFO's, requiring two external memory accesses per video clock as in the case of the *XPGAs* (thus still viable on our *G800* board).

Larger convolution matrices can be considered only using new versions of the *G800* board, featuring faster DRAMs.

A more significant evolution of the present architecture is the implementation of a series of convolutions. In fact, the experiences on the 3×3 matrices showed that both *YPGA* and *XPGA* are underutilized (25% to 30% of CLBs not used).

It seems then feasible to use both *YPGA* and *XPGA* for computation and FIFO management, thus allowing to perform, on the same image, two subsequent convolutions. This can be useful to have an initial image filtering (to enhance signal-to-noise ratio) or to perform edge detection and computation of the brightness gradient (as required, for instance, in the Hough transform).

## References

- [1] D. Buell (editor). *Splash 2 : "FPGA's in a Custom Computing Machine"*. IEEE Computer Society Press, 1996
- [2] Virginio Cantoni, Stefano Levialdi, "*La visione delle macchine – Tecniche e strumenti per l'elaborazione di immagini e il riconoscimento di forme*", Tecniche Nuove, Milano, 1990
- [3] Vito Cappellini, "*Elaborazione numerica delle immagini*", Boringhieri, Torino, 1985
- [4] *The Programmable Logic Data Book*. Xilinx, San Jose, CA, 1996
- [5] Peter Alfke, "*Synchronous and Asynchronous FIFO Designs*". Xilinx Application Note XAPP 051, September 1996, Version 2.0
- [6] Jazi Eko Istiyanto, "*The Formation of Super-Cliques in the Behavioural Synthesis of FIFO RAMs*", Tech Report, Gadjah Mada University
- [7] P.Athanas and L. Abbott. "Addressing the Computational Requirements of Image Processing with a Custom Computing Machine: An Overview". In *Proceedings of the 2<sup>nd</sup> Workshop on Reconfigurable Architecture*, Santa Barbara, CA, April 1995
- [8] Stephen L. Wasson, "*FPGA Design: Early Implications In Partitioning*", Integrated System Design magazine, February 1997
- [9] S. Hauck, *Multi-FPGA Systems* PhD thesis, Department of Computer Science, University of Washington, Sep 1995
- [10] *IEEE Standard VHDL Language Reference Manual*. IEEE Standards
- [11] *FPGA Compiler User Guide v3.5*. Synopsys, Mountain View (USA), Sep 1996
- [12] *Giga Operations Spectrum Documentation*. Giga Operations Corporation, Berkley, CA, 1995